

# GENASIS Basics: Object-oriented utilitarian functionality for large-scale physics simulations (Version 3)

Reuben D. Budiardja<sup>a,1</sup>, Christian Y. Cardall<sup>b,\*</sup>

<sup>a</sup> National Center for Computational Sciences, Oak Ridge National Laboratory, Oak Ridge, TN 37831-6354, USA

<sup>b</sup> Physics Division, Oak Ridge National Laboratory, Oak Ridge, TN 37831-6354, USA

## ARTICLE INFO

### Article history:

Received 12 March 2019

Received in revised form 22 May 2019

Accepted 24 May 2019

Available online 6 June 2019

### Keywords:

Simulation framework

Object-oriented programming

Fortran 2003

## ABSTRACT

GENASIS Basics provides Fortran 2003 classes furnishing extensible object-oriented utilitarian functionality for large-scale physics simulations on distributed memory supercomputers. This functionality includes physical units and constants; display to the screen or standard output device; message passing; I/O to disk; and runtime parameter management and usage statistics. This revision – Version 3 of Basics – includes a significant name change, some minor additions to functionality, and a major addition to functionality: infrastructure facilitating the offloading of computational kernels to devices such as GPUs.

### New version program summary

**Program Title:** SineWaveAdvection, SawtoothWaveAdvection, and RiemannProblem (fluid dynamics example problems illustrating GENASIS Basics); ArgonEquilibrium and ClusterFormation (molecular dynamics example problems illustrating GENASIS Basics)

**Program Files doi:** <http://dx.doi.org/10.17632/6w9ygygmc.2>

**Licensing provisions:** GPLv3

**Programming language:** Fortran 2003 (tested with GNU Compiler Collection 8.1.0, Intel Fortran Compiler 18.0.3, Cray Compiler Environment 8.6.5, IBM XL Fortran 16.1.0)

**Journal reference of previous version:** Computer Physics Communications 214 (2017) 247

**Does the new version supersede the previous version?:** Yes

**Reasons for the new version:** This version includes a significant name change, some minor additions to functionality, and a major addition to functionality: infrastructure facilitating the offloading of computational kernels to devices such as GPUs.

### Summary of revisions:

The class `VariableGroupForm` – a major workhorse for handling set of related fields – has been renamed `StorageForm`.

The ability to use unicode characters in standard output has been added, but is currently only supported by the GNU Compiler Collection (GCC). This capability is used to display exponents as numerical superscripts, as well as symbols such as  $\hbar$ ,  $\odot$ , and  $\text{\AA}$  in the display of relevant units. It is made operational by the line

```
USE_ISO_10646 = 1
```

which is now included in the machine-specific makefile fragments with a GCC suffix in the `Build/Machines` directory.

There are some changes to units and constants. The geometrized units of past releases ( $G = c = k = 1$ , with a fundamental unit of meter) have been replaced by natural units ( $\hbar = c = k = 1$ , with MeV as the fundamental unit). Lorentz–Heaviside electromagnetic units are employed (permeability  $\mu = 1$ ; no factors of  $4\pi$  in the Maxwell equations). This refers to numbers as processed internally by the code; as described in the initial release, users can employ the members of the `UNIT` singleton for input/output purposes, that is, to specify or display numbers with any available units they wish.

\* Corresponding author.

E-mail addresses: [reubendb@ornl.gov](mailto:reubendb@ornl.gov) (R.D. Budiardja), [cardallcy@ornl.gov](mailto:cardallcy@ornl.gov) (C.Y. Cardall).

<sup>1</sup> This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, worldwide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (<http://energy.gov/downloads/doe-public-access-plan>).

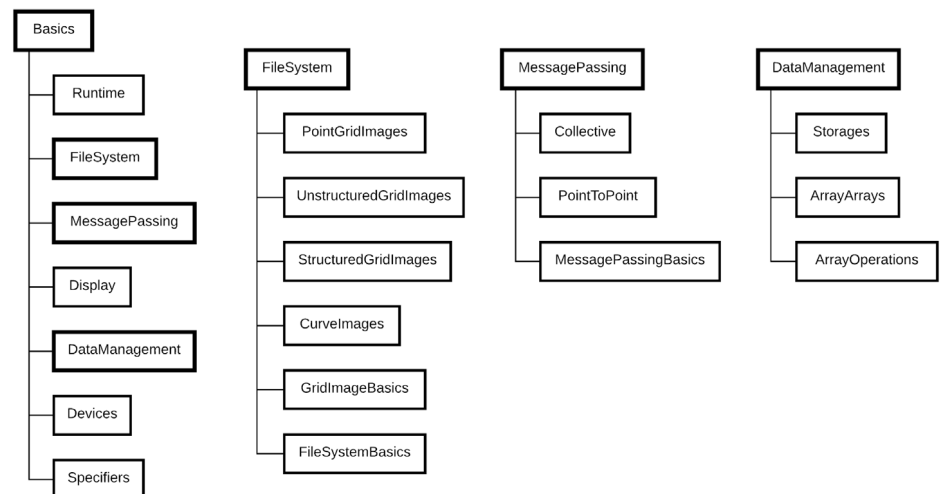
A number of units have been added, and the specification all units has been put on a more rational basis in keeping with six of the seven standard SI base units (meter, kilogram, second, ampere, kelvin, mole; we have not needed the candela; see [3]). Some physical and astrophysical constants have also been added. All constants have been updated to 2018 values [4].

For notifications to standard output, a few tweaks to ignorability levels have been made in various classes. The default output to screen is now less verbose (ignorability INFO\_1, our designation for messages of significance just below WARNING).

A couple of additions have been made to MessagePassing: null subcommunicators are accommodated, and an AllToAll\_V method has been added to CollectiveOperation\_R\_Form.

Enhancements to timer functionality have been made. The class TimerForm now has a member Level, which is specified in order to control indentation in screen output. Some functionality has been added to PROGRAM\_HEADER\_Singleton to work with timers. A method TimerPointer returns a pointer to a timer with a specified Handle (typically a meaningfully named integer). The new members TimerLevel and TimerDisplayFraction of PROGRAM\_HEADER\_Singleton, which can be set from the command line, can be used to suppress output from timings deemed insignificant, based on timer level or a measured time interval falling below a specified fraction of the total execution time.

The most significant addition in functionality in this release is the addition of infrastructure to offload computational kernels to hardware accelerators such as GPUs using OpenMP device-related directives and runtime library routines in OpenMP 4.5 and later.<sup>2</sup> This infrastructure, implemented in a new subdivision Devices (see Fig. 1), provides lower-level routines to perform memory management between the host (CPU) and device (GPU) including data allocation, data movement between host and device, and device-to-host memory address association. The routines are implemented as Fortran wrappers to the OpenMP runtime library and CUDA<sup>3</sup> routines written in C. Additional methods and an option utilizing the lower-level Devices routines have been added to our StorageForm class. They are: UpdateHost() and UpdateDevice() to copy data from device to host and host to device, respectively; AllocateDevice() to allocate memory on the device mirroring the allocation on the host; and PinnedOption as an optional flag to the Initialize() method to allocate the host memory in a page-locked region to facilitate faster data transfer between host and device. A detailed description of the implementation of this functionality can be found in [5].



**Fig. 1.** Left: Structure of Basics after the introduction of the Devices division. Middle Left: Substructure of FileSystem. Middle Right: Substructure of MessagePassing. Right: Substructure of DataManagement. All: Solid lines outline the directory hierarchy. Boxes framed with thinner linewidths denote 'leaf' divisions of the code with no additional subdirectories. The compilation order is from bottom to top; thus dependencies essentially flow in reverse, from top to bottom.

To deal with different levels of compiler support for device-related OpenMP directives, we use the preprocessor in some source files in Devices to guard against attempted compilation of unsupported features. Preprocessor macro substitution is also utilized in OpenMP directives to switch between multi-threading parallelism on CPUs and offload parallelism to GPUs. Setting the makefile variable ENABLE\_OMP\_OFFLOAD to 1 – which is the default in the machine-specific makefile Makefile\_POWER\_XL for the XL compiler on POWER-based supercomputers – sets the appropriate flags and preprocessing to enable compilation for OpenMP offload parallelism. Alternatively, the command

```
make ENABLE_OMP_OFFLOAD=1
```

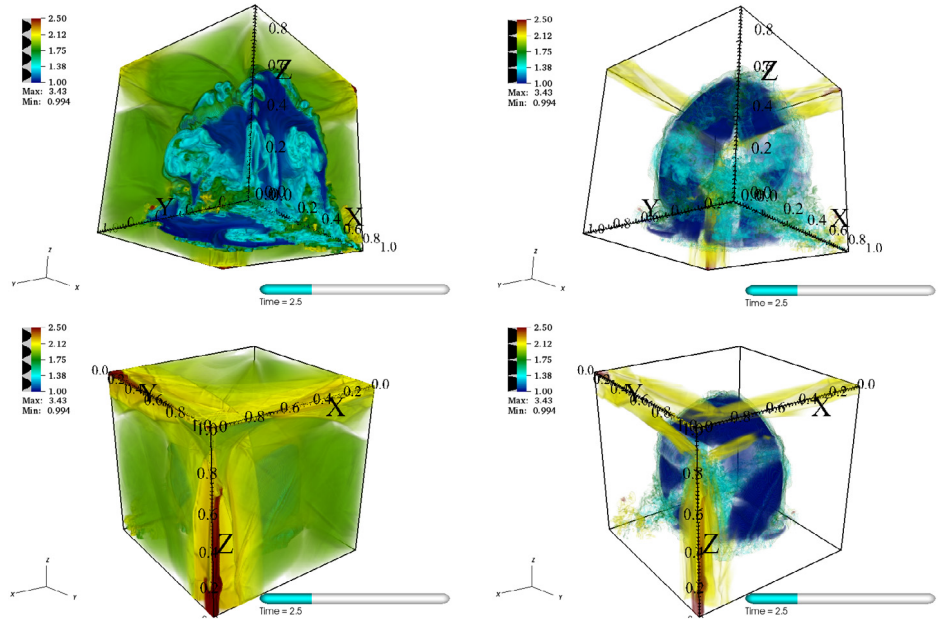
sets this variable when make is invoked from the command line.

<sup>2</sup> <https://www.openmp.org/specifications/>.

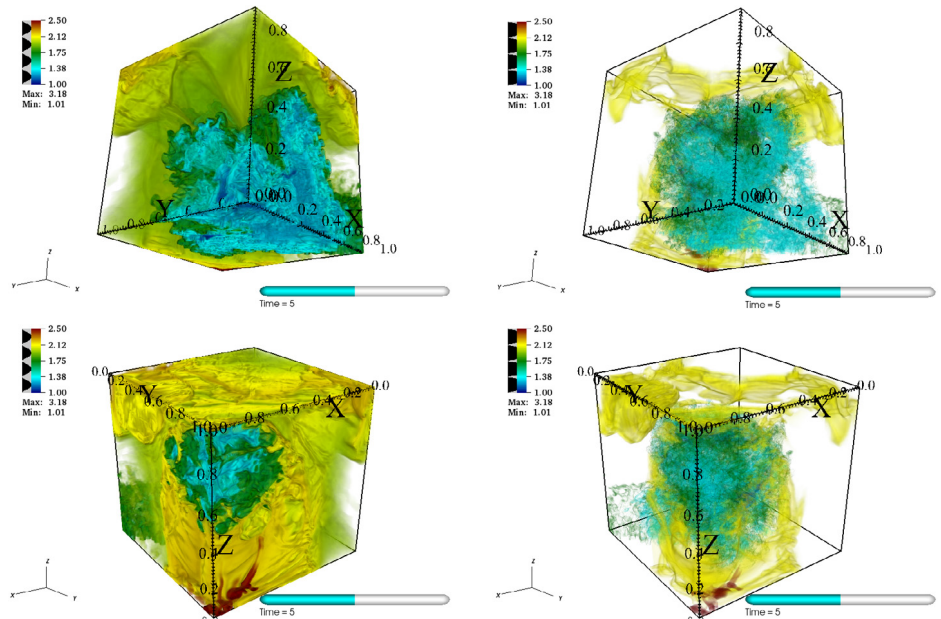
<sup>3</sup> <https://developer.nvidia.com/about-cuda>.

Information regarding the number of devices available to the program, the kind of OpenMP parallelism enabled (i.e. multi-threading or offload), and the selected OpenMP loop scheduling are displayed at runtime by `PROGRAM_HEADER_Singleton`. When offload parallelism is enabled, the loop scheduling is automatically set to `static` with chunk-size of 1. With multi-threading parallelism, the schedule defaults to `guided` but can be overridden at runtime by setting the environment variable `OMP_SCHEDULE` appropriately.

The example problem `RiemannProblem` in the `Examples` directory under the `Basics` division has been modified to exploit the GPUs using this new functionality. The computational kernels for the problem have been annotated with new OpenMP directives (via the appropriate preprocessor macros) such that they are offloaded to the GPUs when offload parallelism is enabled during compilation. In [5] we demonstrate the weak scaling of this example problem up to 8000 GPUs on the Summit supercomputer at the Oak Ridge Leadership Computing Facility.<sup>4</sup> Figs. 2 and 3 show a visualization of the three-dimensional version of `RiemannProblem` at  $1280^3$  resolution executed with 1000 GPUs.



**Fig. 2.** The polytropic constant (essentially, entropy per baryon) in a three-dimensional version of the `RiemannProblem` example at  $1280^3$  resolution at time  $t = 2.5$ . The initial conditions consist of two regions with different pressures and densities separated by a plane of discontinuity intersecting three points:  $[0.5, 0.0, 0.0]$ ,  $[0.0, 0.75, 0.0]$ ,  $[0.0, 0.0, 1.0]$  in a computational domain spanning  $[0, 1]$  in each dimension. Volume plots from two different perspectives (top and bottom) of with two different transfer functions (left and right) are displayed.



**Fig. 3.** As in Fig. 2, but at time  $t = 5$ .

<sup>4</sup> <https://www.olcf.ornl.gov/olcf-resources/compute-systems/summit/>.

*Nature of problem:* By way of illustrating GENASIS Basics functionality, solve example fluid dynamics and molecular dynamics problems.

*Solution method:* For fluid dynamics examples, finite-volume. For molecular dynamics examples, leapfrog and velocity-Verlet integration.

*External routines/libraries:* MPI [1] and Silo [2]

*Additional comments including restrictions and unusual features:* The example problems named above are not ends in themselves, but serve to illustrate our object-oriented approach and the functionality available through GENASIS Basics. In addition to these more substantial examples, we provide individual unit test programs for each of the classes comprised by GENASIS Basics.

GENASIS Basics is available in the CPC Program Library and also at <https://github.com/GenASIS>.

[1] <http://www.mcs.anl.gov/mpi/>

[2] <https://wci.llnl.gov/simulation/computer-codes/silo>

[3] [https://en.wikipedia.org/wiki/SI\\_base\\_unit](https://en.wikipedia.org/wiki/SI_base_unit)

[4] M. Tanabashi et al. (Particle Data Group), Phys. Rev. D 98 (2018) 030001

[5] Budiardja, R.D. and Cardall, C.Y., "Targeting GPUs with OpenMP Directives on Summit: A Simple and Effective Fortran Experience," *submitted for publication Parallel Computing: Systems and Applications*, arXiv:1812.07977 [physics.comp-ph],

© 2019 Elsevier B.V. All rights reserved.

---

## Acknowledgments

This material is based upon work supported by the U.S. Department of Energy, Office of Science, Office of Nuclear Physics

under contract number DE-AC05-00OR22725 and the National Science Foundation under Grant No. 1535130. This research used resources of the Oak Ridge Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC05-00OR22725.